

Table of Contents

1. Qualified References	1
2. Using a Table Alias.....	1
3. SQL Inner Joins	2
4. ANSI inner join: Column Name Join (the USING clause)	3
5. ANSI inner join: Condition Join (the ON clause).....	6
6. NATURAL JOIN.....	9

One of the goals of a relational database is to reduce redundancy; we want to store descriptive data one time only. To do this we need to split data across several tables- so we have a Customer table and an Order table and an OrderDetails table. As a consequence, we will often need to write queries that bring the data back together again from two or more tables. We will need to indicate in the query how these tables are related. We will first discuss table aliases which are commonly used when our query involves more than one table. Then we will look at the ANSI standard inner join syntax.

1. Qualified References

We know that each column in a table needs to have a unique name. But we can have the same identifier for columns in different tables. It is common, and good style, for the pk column in the parent table and the fk column in the child table to have the same name. So if we are joining two tables we may have to qualify any reference to a column name which is the same in the two tables. (The exception is the joining column when we use the Using (col) syntax.) The format for a qualified name is `tblName.ColumnName`.

We can qualify all of the column names in the query. If this is a single table query there is no need to do this. With a multi-table query, the query will be somewhat more efficient if we fully qualify the column names.

For most of the example in this discussion, I will not fully qualify all of the column names since the queries are short and they are easier to read without the qualification.

The purpose of these demos is to show you techniques and syntax- and ease of reading is more important for this purpose than execution efficiency.

For some of our data the tables occurs in different databases. Assuming your account has permissions on all the databases used in the query you can do joins across databases. For these queries the tables will be qualified with the database names.

An ambiguous reference means that you have a column identifier in your query and the same identifier is part of more than one of the tables used in the query. Therefore the system does not know which column is being referenced. This is an error.

2. Using a Table Alias

In MySQL table aliases may be case specific depending on your system. I will try to stick with lower case table aliases in these discussions. My test system is not case specific for table names and table aliases, so if a demo does not work and your system is case specific check for a case issue.

Demo 01: A single table select without table aliases. You can run this query from any database since the From clause specifies the database for the table.

```
select a_prd.warehouses.warehouse_id
, a_prd.warehouses.loc_id
from a_prd.warehouses;
+-----+-----+
| warehouse_id | loc_id |
+-----+-----+
|           100 |    1400 |
|           125 |    1500 |
```

	250	1500
	300	1500
	200	1800

Demo 02: You could also use these versions of the query

```
select warehouses.warehouse_id
, warehouses.loc_id
from a_prd.warehouses;
```

```
select warehouse_id
, loc_id
from a_prd.warehouses;
```

Demo 03: A single table select with a table alias; this will give us the same output as the first query. The table alias is new name to refer to the table.

```
select wh.warehouse_id
, wh.loc_id
from a_prd.warehouses wh;
```

The use as of key word AS is optional with a table alias.

We could rewrite this query as

```
select warehouse_id
, loc_id
from a_prd.warehouses wh;
```

But we cannot use the following; once we set up the alias we cannot use the table name as a qualifier.

```
select a_prd.warehouses.warehouse_id
, a_prd.warehouses.loc_id
from a_prd.warehouses wh;
ERROR 1054 (42S22): Unknown column 'a_prd.warehouses.warehouse_id' in 'field list'
```

3. SQL Inner Joins

We can use several types of joins between tables. We start with inner joins. For a row to appear in the result set returned by an inner join, there needs to be matching data in the joining column(s) in the two tables. There are several ways to implement the inner join.

These examples show the ANSI standard INNER JOIN syntax with the column name joins syntax and the condition join syntax. Most dbms now support this syntax.

The legacy comma join syntax is discussed later. **For assignments in this class you are required to use a syntax that does the join in the From clause.** In a job situation where you are reading and maintaining old code, you will need to understand the legacy syntax which expresses the joining condition as a criterion in the Where clause along with any filter criteria.

We will also talk about the Cartesian product- which is the logical basis for all of the joins.

4. ANSI inner join: Column Name Join (the USING clause)

Let's start with an example of an inner join between the employee table and the department table. We want to see the name of the employee (which is in the employee table) and the name of the department which is in the Department table. If we described this situation to someone we might say that we are **joining** the data in these two tables and that we are **using** the department_id to associate each employee row with the correct department row. We also want to display the department_id. If someone asks us which department_id we want- the one from the employee table or the one from the department table, we would reply that we don't care because they have to be the same value.

The ANSI syntax here models that way of talking about the join.

Demo 04: Inner join employees and their dept

```
select emp_id
, name_last as "Employee"
, dept_name
from a_emp.employees
INNER JOIN a_emp.departments USING (dept_id)
limit 10;
```

emp_id	Employee	dept_name
100	King	Administration
201	Harts	Marketing
101	Koch	Development
108	Green	Development
109	Fiet	Development
110	Chen	Development
203	Mays	Development
204	King	Development
205	Higgs	Development
206	Geitz	Development

Discussion:

With this syntax (USING) , the joining column must have the same name in both tables. This is commonly true of tables in a child-parent situation.

We do not need to qualify the column names in common. We do need to put parentheses around the common column name in the USING clause.

We need to qualify any other column names that would cause an ambiguous reference.

All of the information about the joining of the tables is placed in the From clause. The From clause is supposed to indicate the data source and this syntax defines the data source.

It might seem that the dbms could figure out which columns are the joining columns- particularly if we define the relationships when we create the table- but it doesn't.

Another thing to be aware of: if we have any department with no employees, the inner join will not display that department. To get those rows we would need an outer join which we will discuss soon.

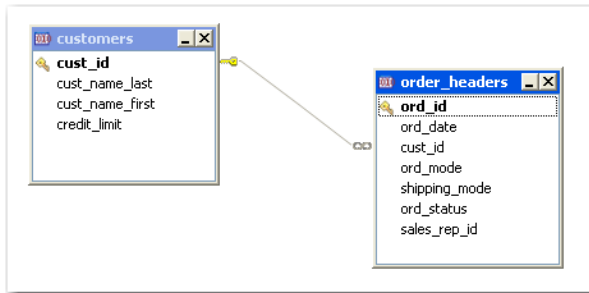
Demo 05: Inner join with USING - show customers and their orders.

```
select cust_id
, cust_name_last as "Customer"
, ord_id
from a_oe.customers
```

```
INNER JOIN a_oe.order_headers USING (cust_id)
order by cust_id
limit 10;
```

cust_id	Customer	ord_id
400300	McGold	378
401250	Morse	106
401250	Morse	113
401250	Morse	119
401250	Morse	301
401250	Morse	506
401890	Northrep	112
401890	Northrep	519
402100	Morise	114
402100	Morise	115

The following is a graphic of a two table join. This is done with the Toad client.



We can join more than two tables- we just add them to the From clause in a logical order and set up the joining column for each pair of tables as we go. In the next query we get one row for each product a customer has ordered. Tables are joined two at a time. The customer table is joined to the orders table to form a virtual table, which is then joined to the order details table to form another virtual table used as the table expression.

Demo 06: three table join; inner join with column name join

```
select cust_id
, cust_name_last as "Customer"
, ord_id
, prod_id
from a_oe.customers
INNER JOIN a_oe.order_headers USING (cust_id)
INNER JOIN a_oe.order_details USING (ord_id)
order by cust_id, ord_id
limit 10;
```

cust_id	Customer	ord_id	prod_id
400300	McGold	378	1125
400300	McGold	378	1120
401250	Morse	106	1060
401250	Morse	113	1080
401250	Morse	119	1070
401250	Morse	301	1100
401890	Northrep	112	1110

	401890		Northrep		519		1020	
	401890		Northrep		519		1110	
	402100		Morise		114		1130	
+-----+-----+-----+-----+								



The keyword INNER is optional and is omitted for the next queries. The default Join is an Inner Join. This style of SQL layout makes it easy to scan down the left edge of the query and see the tables involved.

Demo 07: four table join; inner join with USING


```
select cust_id
, cust_name_last as "Customer"
, ord_id, prod_id, prod_name
from a_oe.customers
JOIN a_oe.order_headers USING (cust_id)
JOIN a_oe.order_details USING (ord_id)
JOIN a_prd.products USING (prod_id)
limit 10;
```

	cust_id		Customer		ord_id		prod_id		prod_name	
+-----+-----+-----+-----+-----+-----+										
	400300		McGold		378		1120		Washer	
	400300		McGold		378		1125		Dryer	
	401250		Morse		106		1060		Mountain bike	
	401250		Morse		113		1080		Cornpopper	
	401250		Morse		119		1070		Iron	
	401250		Morse		301		1100		Blender	
	401890		Northrep		112		1110		Pancake griddle	
	401890		Northrep		519		1020		Dartboard	
	401890		Northrep		519		1110		Pancake griddle	
	402100		Morise		114		1130		Mini Freezer	
+-----+-----+-----+-----+-----+-----+										

Demo 08: five table join including a row filter for appliances .

```
select cust_id
, ord_id
, prod_id
, prod_name
, quoted_price
from a_oe.customers
join a_oe.order_headers using (cust_id)
join a_oe.order_details using (ord_id)
join a_prd.products using (prod_id)
join a_prd.categories using (catg_id)
where catg_desc in( 'APPLIANCES' )
limit 10;
```

cust_id	ord_id	prod_id	prod_name	quoted_price
402100	115	1120	Washer	475.00
409030	130	1120	Washer	500.00
903000	306	1120	Washer	500.00
900300	307	1120	Washer	450.00
400300	378	1120	Washer	450.00
403010	118	1125	Dryer	475.00
409030	130	1125	Dryer	500.00
903000	306	1125	Dryer	500.00
900300	307	1125	Dryer	450.00
400300	378	1125	Dryer	450.00



5. ANSI inner join: Condition Join (the ON clause)

Sometimes we need to join two tables that have different names for the joining column. In this example we are joining the order table to the employee table. In the employee table employees are identified with an employee_id. The order table uses the term sales_rep_id to refer to the same values.

Now we use the key word ON instead of Using and we list the two columns with an equality operator. We would not really have to qualify these columns but it is common, and helpful, to do so.

Demo 09: two table join- inner join with ON clause. In this case, the joining columns have different names and you cannot have a USING clause.

It is not required to use table aliases -but it does make the query easier to read.

```

select ord_id
, cust_id
, emp_id
, name_last as "SalesRep"
from a_oe.order_headers oh
join a_emp.employees      em on oh.sales_rep_id = em.emp_id
limit 10;

```

ord_id	cust_id	emp_id	SalesRep
105	403000	150	Tuck
106	401250	150	Tuck
107	403050	150	Tuck
108	403000	155	Hiller
109	403000	155	Hiller
110	404950	155	Hiller
111	403000	150	Tuck
112	401890	145	Russ
113	401250	150	Tuck
114	402100	155	Hiller

It is legal to use the On syntax when you are joining two tables which use the same column name for the joining column.

Demo 10: Inner join with ON clause - Show customers and their orders. Now you have to qualify the cust_id column because we are not using the column name join.

```
select cs.cust_id
, cust_name_last as "Customer"
, ord_id
from a_oe.customers      cs
join a_oe.order_headers oh on cs.cust_id = oh.cust_id
limit 10;
```

cust_id	Customer	ord_id
400300	McGold	378
401250	Morse	106
401250	Morse	113
401250	Morse	119
401250	Morse	301
401250	Morse	506
401890	Northrep	112
401890	Northrep	519
402100	Morise	114
402100	Morise	115

Demo 11: The only column in the select you are required to qualify is cust_id because it occurs in two tables in the From clause. But it is common and good style to qualify all of the columns when you have a join. This is easier to read with table aliases.

```
select cs.cust_id
, cs.cust_name_last as "Customer"
, oh.ord_id
from a_oe.customers      cs
join a_oe.order_headers oh on cs.cust_id = oh.cust_id
limit 10;
```

You can combine both syntaxes in the same query.

Demo 12: Inner join with a USING and an ON clause- Show customers and their orders and their sales rep.

```
select cs.cust_id
, cs.cust_name_last as "Customer"
, oh.ord_id
, em.emp_id
, em.name_last as "SalesRep"
from a_oe.customers      cs
join a_oe.order_headers oh using (cust_id)
join a_emp.employees     em on oh.sales_rep_id = em.emp_id
limit 10;
```

cust_id	Customer	ord_id	emp_id	SalesRep
400300	McGold	378	150	Tuck
401250	Morse	106	150	Tuck
401250	Morse	113	150	Tuck

	401250		Morse		119		155		Hiller	
	401250		Morse		301		150		Tuck	
	401250		Morse		506		150		Tuck	
	401890		Northrep		112		145		Russ	
	401890		Northrep		519		155		Hiller	
	402100		Morise		114		155		Hiller	
	402100		Morise		117		150		Tuck	
+-----+-----+-----+-----+-----+										

Demo 13: This is the five table join including a row filter using the condition join syntax.

I find it easier to get all the joins written if I have a pattern. My pattern is to use 2 character table aliases. This reduces the problems with using O for the order_headers table and then using D for the order details table.

When I use the condition join syntax, I write the join as the prior table.col = this table.col..

```
select
  cs.cust_id
, oh.ord_id
, OD.prod_id
, PR.prod_name
, OD.quoted_price
from a_oe.customers      cs
join a_oe.order_headers  oh  on cs.cust_id = oh.Cust_id
join a_oe.order_details  od  on oh.ord_id = od.ord_id
join a_prd.products      pr  on od.prod_id = pr.prod_id
join a_prd.categories    ct  on pr.catg_id = ct.catg_id
where ct.catg_desc in( 'APPLIANCES' )
limit 10;
```

Demo 14: This is a style of query I often get from students. This is barely readable and has an error. How easy is it to see the error- without reformatting the query?

```
select a_oe.customers.cust_id, a_oe.order_headers.ord_id,
a_oe.order_details.prod_id,
a_prd.products.prod_name,a_oe.order_details.quoted_price
from a_oe.customers join a_oe.order_headers on a_oe.customers.cust_id =
a_oe.order_headers.Cust_id join a_oe.order_details on
a_oe.order_details.ord_id = a_oe.order_headers.ord_id
join a_prd.products on a_oe.order_details.prod_id =
a_oe.order_details.prod_id join a_prd.categories on a_prd.categories.catg_id
= a_prd.products.catg_id
where a_prd.categories.catg_desc in( 'APPLIANCES' )
limit 10;
```

There is a small difference in using a Select * from tbl1 join tbl2 depending on whether you use the column name join and the condition join. The condition join show two copies on the joining column and the column name join shows only one column in the result set with the joining column.

Demo 15:

```
select *
from a_prd.warehouses w
join a_prd.inventory i using(warehouse_id);
+-----+-----+-----+-----+-----+
| warehouse_id | loc_id | prod_id | quantity_on_hand |
```


100	1400	1000	80	
100	1400	1010	500	

```
select *
from a_prd.warehouses w
join a_prd.inventory i on w.warehouse_id = i.warehouse_id;
```

warehouse_id	loc_id	prod_id	warehouse_id	quantity_on_hand
100	1400	1000	100	80
100	1400	1010	100	500

6. NATURAL JOIN

This join can be used if the columns in common have the same name. This will join on any and all columns having the same identifier in the two tables. Therefore it creates maintenance problems if attributes are renamed or additional columns are added to the tables. Suppose we had two tables such as a customer table and a salesrep table which should be joined on a salesrepID column. But perhaps we also have attributes named City and State in each of these tables. If we did a natural join, the join would be on all of these columns. We will not use the Natural Join syntax in this class.

If you specify a natural join on tables that do not have any column names in common, then you get a cross join.